GEORG-AUGUST-UNIVERSITÄT GÖTTINGEN IN PUBLICA COMMODA SEIT 1737

Project Thesis

submitted in partial fulfillment of the
requirements for the course "Applied Computer Science"

A comparision of various deep learning based detection systems for multi-object tracking

Lennart Hahner

MatrNr: 16853416

Supervisor: Lukas Steinegger

Georg-August-Universität Göttingen
Institute of Computer Science

ISSN: 1612-6793

February 25, 2026

Georg-August-Universität Göttingen
Institute of Computer Science

Goldschmidtstraße 7
37077 Göttingen
Germany

☎ +49 (551) 39-172000
FAX +49 (551) 39-14403
✉ office@informatik.uni-goettingen.de
www.informatik.uni-goettingen.de

Abstract

[Put your abstract here.]

General structure of an abstract, write 1 to 2 sentences per section

1. A general statement introducing the broad research area of the particular topic being investigated.
2. An explanation of the specific problem (difficulty, obstacle, challenge) to be solved.
3. A review of existing or standard solutions to this problem and their limitations.
4. An outline of the proposed new solution.
5. A summary of how the solution was evaluated and what the outcomes of the evaluation were.

You may find the following resources useful:

- `https://www.grammarly.com/blog/write-an-abstract/`
- `https://www.editage.com/insights/manuscript-structure-how-to-convey-your-most-im`
- More useful links: `https://hps.vi4io.org/teaching/ressources/start`

Declaration on the use of ChatGPT and comparable tools in the context of examinations

In this work I have used ChatGPT or another AI as follows:

- ☑ Not at all
- ☐ During brainstorming
- ☐ When creating the outline
- ☐ To write individual passages, altogether to the extent of 0% of the entire text
- ☐ For the development of software source texts
- ☐ For optimizing or restructuring software source texts
- ☐ For proofreading or optimizing
- ☐ Further, namely: -

I hereby declare that I have stated all uses completely.
Missing or incorrect information will be considered as an attempt to cheat.

Contents

List of Tables **iv**

List of Figures **v**

List of Listings **vi**

List of Abbreviations **vii**

1 Introduction **1**

2 Related work **1**
2.1 Overview of modern detection systems 1
2.2 Integration of detection systems in tracking systems 2

3 Experimental setting **4**
3.1 Datasets 4
3.1.1 MOT15 Dataset collection 4
3.1.2 MOT20 Dataset collection 5
3.2 Detection systems 6
3.3 Multi-object-tracking systems 6

4 Benchmark results **7**
4.1 Metrics 7
4.2 Analysis 7

5 Discussion **8**

References **9**

A Additional Materials **A1**

List of Tables

List of Figures

1 Object detection systems over time 2
2 4
3 5

List of Listings

List of Abbreviations

HPC High-Performance Computing

1 Introduction

2 Related work

2.1 Overview of modern detection systems

Edozie et al. [1] provide a comprehensive review of the development of object detection techniques over the past 26 years. In addition to describing the evolution of deep learning–based object detection systems, they compare these approaches with classical object detection methods that do not rely on deep learning.

Their survey presents a broad overview of relevant datasets, methods, and research contributions over this period. In addition to classical detection methods, they also review modern deep learning–based object detectors, providing a clear overview of the current state of the art.

In their work, they propose several ways to categorize object detection systems. One common categorization distinguishes between one-stage and two-stage detectors. Two-stage detectors first generate region proposals and subsequently classify these proposals into object categories. In contrast, one-stage detectors do not explicitly generate region proposals but instead directly predict object classes and bounding boxes in a single step. Due to this simplified processing pipeline, one-stage detectors generally achieve higher inference speeds than two-stage detectors.

However, this categorization does not explicitly account for newer architectures such as transformer-based models, e.g., DETR <**empty citation**> or multi-path convolutional neural networks such as PANet <**empty citation**> Therefore, the authors additionally classify detection architectures into four categories: traditional detectors, which include non–deep-learning-based methods; conventional deep learning detectors, primarily consisting of two-stage detection systems; hybrid or bridging methods, which include multi-path detection systems; and advanced deep learning methods, which include one-stage detectors and transformer-based models. Figure 1 illustrates an overview of different object detectors, their respective categories, and their year of development, as presented in [1].

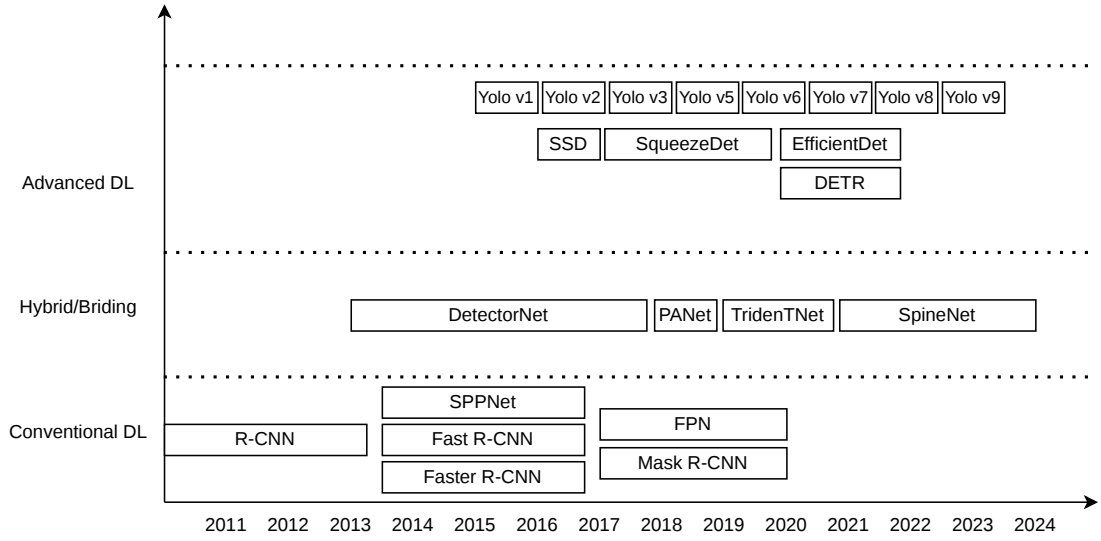

Figure 1: This graphic summarizes the development of object detection systems over recent years, including R-CNN <**empty citation**> SPP-Net <**empty citation**> Fast R-CNN <**empty citation**> Faster R-CNN <**empty citation**> FPN <**empty citation**> Mask R-CNN <**empty citation**> DetectoRS <**empty citation**> PANet <**empty citation**> TridentNet <**empty citation**> SpineNet <**empty citation**> SSD <**empty citation**> SqueezeDet <**empty citation**> EfficientDet <**empty citation**> and DETR <**empty citation**>

Other frameworks, such as open-vocabulary detection and 3D object detection, are not included in the figure but are discussed and categorized in [1].

Similarly, Zou et al. [2] present a comprehensive survey of object detection methods. They also distinguish between one-stage and two-stage detectors and provide additional insights into modern transformer-based detection architectures. Furthermore, they compare different detection methods, as illustrated in Figure 1, based on their historical development, architectural design, and performance. Performance is evaluated using the mean Average Precision (mAP) metric on benchmark datasets such as VOC07 and MS COCO. Their results show that modern detection systems achieve mAP values exceeding 70%.

Neha et al. [3] provide an additional overview of object detection systems based on both classical and deep learning approaches. They summarize relevant datasets, evaluation metrics, and detection algorithms, and discuss the strengths and limitations of various deep learning–based detectors. Furthermore, they present benchmark results for representative detection models such as R-CNN and YOLO.

2.2 Integration of detection systems in tracking systems

Li et al. **li** provide a comprehensive overview of modern multi-object tracking (MOT) systems and their integration with object detection methods. They categorize MOT approaches into three main groups: detection-based methods, single-object tracking–based methods, and segmentation-based methods. Detection-based methods first detect objects using an object detector and subsequently associate these detections across frames us-

ing motion and appearance features. In contrast, single-object tracking–based methods use individual object trackers to follow multiple objects across frames. Segmentation-based methods, on the other hand, track objects using pixel-level segmentation instead of bounding boxes, enabling more precise object localization.

The authors emphasize that the performance of multi-object tracking systems strongly depends on the accuracy of the underlying object detector. Modern tracking systems commonly use detectors such as YOLO, CenterNet, and DETR. Their survey also compares multiple tracking methods and shows that modern approaches, such as ByteTrack and BoT-SORT, achieve MOTA values exceeding 80% on the MOT17 dataset. Both ByteTrack and BoT-SORT rely on YOLO-based object detection models.

Adžemović **adzemovic** presents a similar categorization of tracking methods and introduces an additional paradigm referred to as end-to-end tracking, in which detection and tracking are learned jointly within a single neural network. This category includes transformer-based tracking methods such as TrackFormer and MOTR, which jointly learn object detection and association. The survey highlights the performance advantages of end-to-end tracking systems, particularly MOTRv2 <**empty citation**> which achieves state-of-the-art tracking performance. In contrast, ByteTrack achieves the best MOTA performance among tracking-by-detection methods.

Bashar et al. **bashar** also identify tracking-by-detection as the dominant paradigm in modern multi-object tracking systems, where objects are first detected and then associated across frames. Recent approaches increasingly employ transformer-based architectures, such as TrackFormer and TransTrack, which improve tracking performance by modeling temporal relationships between frames. Additionally, graph neural networks, motion models, and Siamese networks are used to enhance object association and identity preservation. The survey reports that modern deep learning–based tracking methods achieve MOTA values of up to 78.3% on benchmark datasets such as MOT17 and MOT20, demonstrating significant improvements in tracking performance.

3 Experimental setting

3.1 Datasets

The choice of dataset is important to correctly reflect the actual performance of different permutations of detection and tracking systems. Especially newer datasets that include night-time scenes, non-urban environments, or highly occluded images can provide useful insights into performance improvements over time. Furthermore, this section highlights the evolution of various dataset collections and challenges.

3.1.1 MOT15 Dataset collection

Different permutations of tracking and detection systems were evaluated on the MOT15 dataset [4]. The MOT15 challenge is a dataset framework that contains 5503 training frames and 5783 testing frames, divided into a total of 22 sequences. The dataset includes both 3D and 2D sequences; however, for the current benchmark, only the 2D sequences are considered.

Furthermore, the dataset provides a diverse set of scenes and categorizes the data according to camera movement (static or moving), viewpoint height, and weather conditions. The ground truth annotations are primarily generated through manual human labeling, supported by annotation tools such as the VATIC annotation tool. Some sequences originate from existing datasets, such as the ETH <**empty citation**> dataset and the KITTI <**empty citation**> dataset, which provide their own annotated data. In Figure 2, exemplary ground truth annotations are presented.

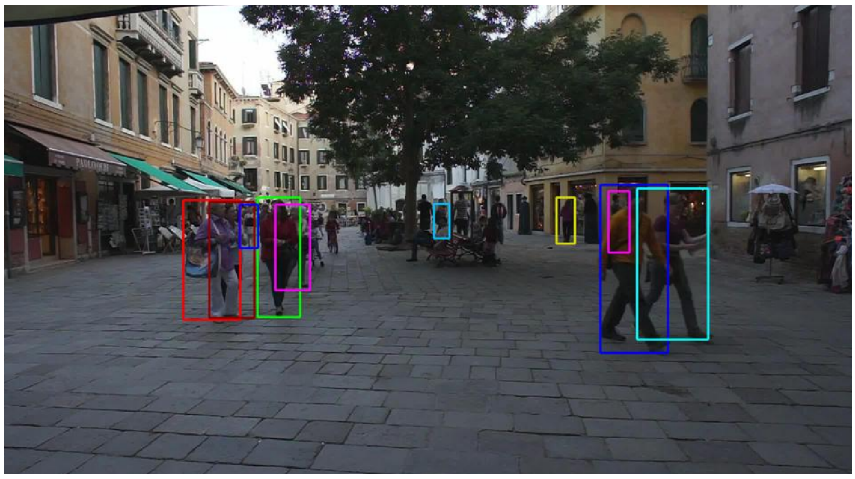

Figure 2: Example scene from the Venice-2 sequence, which is included in the MOT15 dataset. The frames show the point of view (POV) of a static camera in the inner city of Venice. The bounding boxes represent the ground truth annotations provided by the MOT15 challenge. [4]

This figure shows one of the targets used in the benchmark. Most sequences, such as KITTI-17, TUD-Crossing, and ETH-Crossing, depict urban environments. The dataset mainly consists of scenes in which multiple pedestrians are visible. As a result, occlusion is common due to overlapping pedestrians. Many sequences originate from datasets such as ETH-Bahnhof, ETH-Sunnyday, and ETH-Pedcross2.

As previously mentioned, the dataset is divided into two parts: the training set and the testing set. While the training and testing sets contain data from similar environments, they are recorded at different times. For example, the KITTI-17 sequence is from 2017 and is located in the training set, while the KITTI-19 sequence is from 2019 and is located in the testing set. This structure makes it possible to train and evaluate detection and tracking models under similar but temporally distinct conditions. Since this report focuses on pre-trained models, the evaluation is independent of whether the training or testing dataset is used.

3.1.2 MOT20 Dataset collection

Since the MOT15 dataset primarily includes daytime scenes with moderate crowd density, an additional dataset from a newer MOTChallenge benchmark is considered. The MOT20 dataset [5] includes scenes recorded during both daytime and nighttime, and in significantly more crowded environments. This provides additional challenges for detection and tracking systems and allows evaluation under more difficult and realistic real-world conditions.

The MOT20 dataset consists of a total of 8 sequences, split into 4 training and 4 testing sequences. The training set contains 8,931 frames, while the test set contains 4,479 frames. Compared to the MOT15 challenge, the pedestrian density per image has increased substantially, reaching peak densities of up to 246 pedestrians per frame. Furthermore, the dataset contains approximately three times more bounding box annotations than the MOT17 dataset, significantly increasing the complexity of the tracking task. This increase in annotation density and scene complexity is visible in the ground truth example shown in Figure 3.

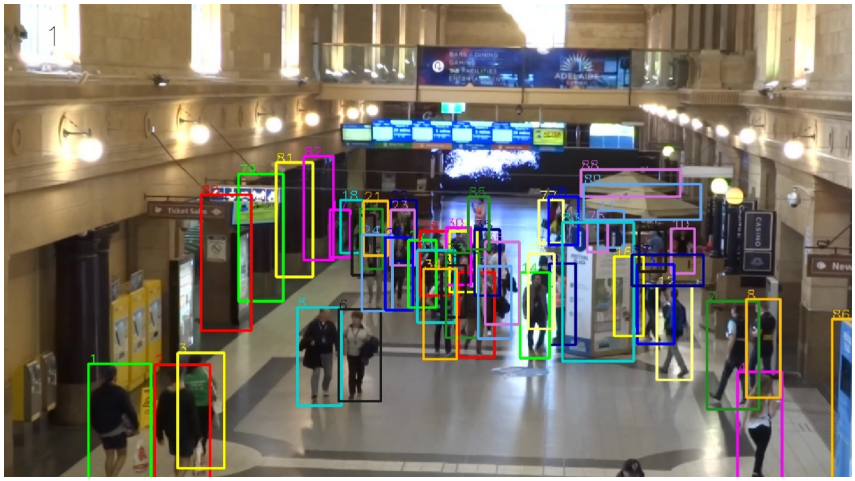

Figure 3: Example scene from the MOT20 dataset. The frames show the point of view of a static camera in a crowded indoor public environment. The bounding boxes represent the ground truth pedestrian annotations provided by the MOT20 challenge. The high number of overlapping annotations illustrates the increased scene complexity and crowd density. [5]

The high pedestrian density introduces additional challenges for MOT systems, as frequent occlusions between individuals occur. This makes it more difficult to maintain correct identity assignment and to associate object tracks consistently across multiple frames. As illustrated in Figure 3, the dataset includes both indoor and outdoor scenes,

recorded from elevated static camera viewpoints. These scenes contain a wide range of environmental conditions, including daytime and nighttime recordings, varying illumination conditions, and different crowd dynamics.

Additionally, all sequences are recorded in high resolution, typically up to 1920×1080 pixels, allowing precise spatial localization of pedestrians. The dataset provides manually annotated ground truth bounding boxes, including unique object identities and visibility information. This enables a detailed evaluation of tracking performance, particularly in challenging scenarios with partial and full occlusions.

3.2 Detection systems

3.3 Multi-object-tracking systems

4 Benchmark results

4.1 Metrics

4.2 Analysis

5 Discussion

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

References

[1] E. Edozie, A. N. Shuaibu, U. K. John, *et al.*, "Comprehensive review of recent developments in visual object detection based on deep learning," *Artif Intell Rev*, vol. 58, p. 277, Jun. 2025. DOI: 10.1007/s10462-025-11284-w. [Online]. Available: https://doi.org/10.1007/s10462-025-11284-w.

[2] Z. Zou, K. Chen, Z. Shi, Y. Guo, and J. Ye, "Object detection in 20 years: A survey," *Proceedings of the IEEE*, vol. 111, no. 3, pp. 257–276, 2023. DOI: 10.1109/JPROC.2023.3238524.

[3] F. Neha, D. Bhati, D. K. Shukla, and M. Amiruzzaman, *From classical techniques to convolution-based models: A review of object detection algorithms*, 2024. arXiv: 2412.05252 [cs.CV]. [Online]. Available: https://arxiv.org/abs/2412.05252.

[4] L. Leal-Taixé, A. Milan, I. Reid, S. Roth, and K. Schindler, "MOTChallenge 2015: Towards a benchmark for multi-target tracking," *arXiv:1504.01942 [cs]*, Apr. 2015, arXiv: 1504.01942. [Online]. Available: http://arxiv.org/abs/1504.01942.

[5] P. Dendorfer, H. Rezatofighi, A. Milan, *et al.*, *Mot20: A benchmark for multi object tracking in crowded scenes*, 2020. arXiv: 2003.09003 [cs.CV]. [Online]. Available: https://arxiv.org/abs/2003.09003.

A Additional Materials